

The article explains ValuJet Flight 592's 1996 crash and argues that its root cause is best understood not as a single issue or an individual engineering flaw but as a system accident—what Charles Perrow calls a “normal accident” in tightly coupled, interactively complex systems. The narrative begins with eyewitness accounts and the rapid response of air traffic control, then follows the National Transportation Safety Board's (NTSB) investigation as it focuses on the airplane's forward cargo hold, which lacked fire detection/suppression systems. In there was a lethal mix of issues: chemical oxygen generators (removed from other aircraft during maintenance), cardboard boxes, and aircraft tires—ingredients that, once an oxygen generator fired, produced intense heat and a blowtorch effect.

Crucially, no single issue explains the disaster. Instead, the article maps how organizational complexity and economic pressure—a deregulated, cost-cutting airline (ValuJet), heavy outsourcing to a maintenance contractor (SabreTech), layers of temporary labor, and an FAA attempting to supervise through paperwork—combined into a brittle system. The oxygen generators were removed without the required safety caps; ambiguous documentation (engineer-speak such as “expired” vs. “expended”) and hurried sign-offs normalized shortcuts; unmarked boxes were shipped as routine “company material”; a ramp agent and the copilot—both trained to spot hazardous cargo—failed to intervene. Each link seems minor and understandable in isolation; aligned together under time pressure, they formed a catastrophic chain.

The article distinguishes three types of accidents: procedural (avoidable rule violations), engineered (design/material failures that can be ultimately solved through better engineering), and system accidents (failures that emerge from the functioning of the entire

system—organization, oversight, incentives, language, and culture). The NTSB and press ultimately prompted near-term corrective actions (banning oxygen generators as cargo on passenger flights; promises to retrofit cargo holds), and the FAA leadership changed. But the author warns that piling on more rules and safety devices can increase opacity and coupling that also has the potential to add even more risk. The broader conclusion is simple: aviation is very safe and getting safer, but in complex systems, some accidents are inherently “normal,” and the dream of a zero-accident future is unrealistic. The practical lesson is to avoid building pretend realities—paper processes and language that drift away from the messy constraints of actual work.

Systems of systems in software. Modern software is rarely a single codebase. Even a small 2D game like mine is a system of systems: engine/runtime (e.g., Unity/Godot), platform SDKs (Windows/macOS/iOS/Android), physics/audio/input subsystems, third-party packages (analytics, ads, crash reporting, in-app purchases), build/CI, asset pipelines (art, audio, localization), cloud backends (leaderboards, save sync), and distribution channels (Steam, itch.io, app stores). Each subsystem evolves on its own schedule and exposes ambiguous interfaces, docs, and constraints. Tight coupling shows up as frame-timing dependencies, physics determinism tied to time step, or build scripts that assume specific plugin versions. Interactive complexity is everything that can go “slightly wrong” but usually doesn’t—until several small things line up.

The ValuJet story is a case study in how teams become comfortable with “temporary” exceptions—missing safety caps, unmarked boxes, habitual sign-offs. Games have a parallel: bypassing code review “just this once,” checking in a debug flag left on, committing assets that exceed memory budgets, hard-wiring a URL, or silencing flaky unit tests. Because the build

ships, nothing “blows up,” and the shortcut becomes tacit policy. Over time, the project accumulates fragile assumptions (“we always build in this order,” “rename this folder and the importer breaks”). The danger appears only when deadlines, platform updates, and hotfixes show up together

Engineer-speak and ambiguous requirements. “Expired” vs. “expended” parallels eerily well to “deprecated” vs. “disabled,” “preview” vs. “production-ready,” “editor only” vs. “runtime safe,” “thread-safe” vs. “re-entrant,” etc. If my asset pipeline says “do not compress” in one doc and “must compress” in another, artists and integrators will guess. In game development, language precision matters. Write procedures for the audience that will use them (designers, artists, contractors), not for the engineers who wrote them.

There is a parallel to “safety devices that add risk.” In the crash, oxygen generators—meant for emergencies—became fuel under the wrong conditions. In software, extra layers can add a failure surface: an anti-cheat driver that blue-screens, a crash handler that deadlocks, a watchdog that kills the main thread, or a retry loop that inadvertently amplifies load. Good intentions can create tight coupling and opaque failure modes. For my game, I should prefer simple, observable safeguards: timeouts with backoff, idempotent writes to storage, feature flags that truly bypass code paths rather than merely hiding UI, and circuit breakers around network calls so offline play remains stable.

Designing for slack helps absorb surprises. I want the game loop to remain independent from I/O so the frame rate is resilient to network stalls. Queues and background workers let analytics and cloud saves proceed without touching the render path. Similarly, I should decouple asset import from runtime loading: validate assets on import (size, format, pivot, compression)

and generate runtime-ready bundles rather than “fixing it at runtime.” These design choices reduce tight coupling and give me time to respond when one subsystem misbehaves.